





Namespaces.go

2018




What is a container?


What is a container?



What is a container?



What is a container?






What is a container?

```
$ ps ax
```

PID	TTY	STAT	TIME	COMMAND
1	?	Ss	4:19	/sbin/init splash
2	?	S	0:00	[kthreadd]
3	?	I<	0:00	[rcu_gp]
4	?	I<	0:00	[rcu_par_gp]
6	?	I<	0:00	[kworker/0:0H-kblockd]
8	?	I<	0:00	[mm_percpu_wq]
9	?	S	0:25	[ksoftirqd/0]
10	?	I	3:54	[rcu_sched]
11	?	I	0:00	[rcu_bh]
12	?	S	0:00	[migration/0]

Is container just a feeling? 🍷



Abused containers?

Rewind: 2002-2006

name

—

namespace



What is a namespace?

- Partitions created by Linux kernel
- Isolate views of resources for Processes.
- Every Process has a namespace.
- Init starts with a default namespace.
- Two processes under same namespace have same access to resources.



Where are they?

```
$ ls -al /proc/10504/ns/
```

```
lrwxrwxrwx 1 cgroup -> cgroup:[4026531835]  
lrwxrwxrwx 1 ipc -> ipc:[4026531839]  
lrwxrwxrwx 1 mnt -> mnt:[4026531840]  
lrwxrwxrwx 1 net -> net:[4026532009]  
lrwxrwxrwx 1 pid -> pid:[4026531836]  
lrwxrwxrwx 1 user -> user:[4026531837]  
lrwxrwxrwx 1 uts -> uts:[4026531838]
```

```
$ ls -al /proc/3330/ns/
```

```
lrwxrwxrwx 1 cgroup -> cgroup:[4026531835]  
lrwxrwxrwx 1 ipc -> ipc:[4026531839]  
lrwxrwxrwx 1 mnt -> mnt:[4026531840]  
lrwxrwxrwx 1 net -> net:[4026532009]  
lrwxrwxrwx 1 pid -> pid:[4026531836]  
lrwxrwxrwx 1 user -> user:[4026531837]  
lrwxrwxrwx 1 uts -> uts:[4026531838]
```



Types of namespaces

- Mount
- UTS
- IPC
- PID
- Network
- User
- Cgroup



Containers?

- Containers are to Process what Process are to Threads.
- Virtualization
- Less sharing
- More Separation

**Sharing is not Caring,
Your parents have been lying.**



Syscall?


```
func main() {  
    pid, _, _ := syscall.Syscall(syscall.SYS_GETPID, 0, 0, 0)  
    fmt.Println("process id: ", pid)  
}
```



Syscall

```
$ go run hello.go  
process id: 12566
```

Skeleton





framework

```
func main() {  
    cmd := exec.Command("/bin/sh")  
    cmd.Run()  
}
```

```
$ go run hello.go  
$
```



framework

```
func main() {  
    cmd := exec.Command("/bin/bash")  
  
    cmd.Stdin = os.Stdin  
    cmd.Stdout = os.Stdout  
    cmd.Stderr = os.Stderr  
  
    cmd.Run()  
}
```



framework

```
$ sudo ./main
```

```
root@xps:~/workspace/meson10/test# exit
```



Helper functions

```
// Attaches stdin, stdout, stderr to Cmd.
```

```
func makeCmd(cmd *exec.Cmd) *exec.Cmd {
```

```
    cmd.Stdin = os.Stdin
```

```
    cmd.Stdout = os.Stdout
```

```
    cmd.Stderr = os.Stderr
```

```
    return cmd
```

```
}
```

UTS namespace



UTS namespace

- Isolate uname system call.
- hostname
- domainname

UTS namespace

2 Note how the system call is `uname` but the structure it returns is called `utsname`. In that sense, it seems you can pretty much read UTS == UNIX. Presumably it's called a "UTS namespace", since that hints at `uname`, rather than "UNIX namespace", which suggests something that affects the whole system. – Mikel Feb 9 '15 at 4:50

Timesharing? Why doesn't it have separate system date and time? It would be useful for starting a program that only works in a limited range of dates. – Vi. Feb 9 '15 at 10:13

1 @Vi. That's not what timesharing means. – immibis Oct 10 '16 at 22:29



UTS namespace

```
func main() {  
    cmd := makeCmd(exec.Command("/bin/sh"))  
    cmd.SysProcAttr = &syscall.SysProcAttr{  
        Cloneflags: syscall.CLONE_NEWUTS,  
    }  
    cmd.Run()  
}
```




UTS namespace

```
$ sudo ./main  
# hostname hello  
# hostname  
Hello
```

ANOTHER SHELL

```
$ hostname  
xps.piyushverma.net
```

Picture or it didn't happen





\$: sudo go run hello.go

20078	pts/1	S	0:00				_ sudo ./hello
20079	pts/1	Sl	0:00				_ ./hello
20084	pts/1	S+	0:00				_ /bin/bash

—

User namespace



User namespace

```
func main() {  
    cmd := makeCmd(exec.Command("/bin/sh"))  
    cmd.Env = []string{"PS1=[gophercon]$ "  
  
    cmd.SysProcAttr = &syscall.SysProcAttr{  
        Cloneflags: syscall.CLONE_NEWUSER,  
    }  
    cmd.Run()  
}
```



UID/GID

```
$ echo $UID
```


```
1000
```

```
$ go run user_ns.go
```

```
[gophercon]$ echo $UID
```

```
[gophercon]$
```

Picture or it didn't happen





User namespace

```
$ ls -al /proc/5054/ns/
```

```
lrwxrwxrwx cgroup -> 'cgroup:[4026531835]'  
lrwxrwxrwx ipc -> 'ipc:[4026531839]'  
lrwxrwxrwx mnt -> 'mnt:[4026532471]'  
lrwxrwxrwx net -> 'net:[4026532008]'  
lrwxrwxrwx pid -> 'pid:[4026531836]'  
lrwxrwxrwx user -> 'user:[4026531837]'  
lrwxrwxrwx uts -> 'uts:[4026531838]'
```

```
$ ls -l /proc/3692/ns/
```

```
lrwxrwxrwx cgroup -> 'cgroup:[4026531835]'  
lrwxrwxrwx ipc -> 'ipc:[4026531839]'  
lrwxrwxrwx mnt -> 'mnt:[4026532471]'  
lrwxrwxrwx net -> 'net:[4026532008]'  
lrwxrwxrwx pid -> 'pid:[4026531836]'  
lrwxrwxrwx user -> 'user:[4026532334]'  
lrwxrwxrwx uts -> 'uts:[4026531838]'
```



User namespace

```
$ go run user_ns.go
```

```
[gophercon]$ whoami  
nobody
```



UID/GID

```
cmd := makeCmd(exec.Command("/bin/sh"))
cmd.SysProcAttr = &syscall.SysProcAttr{
    Cloneflags: syscall.CLONE_NEWUSER,
    UidMappings: []syscall.SysProcIDMap{{
        ContainerID: 109,
        HostID:      os.Getuid(),
        Size:        1,
    }},
    GidMappings: []syscall.SysProcIDMap{{
        ContainerID: 114,
        HostID:      os.Getgid(),
        Size:        1,
    }},
}
```



UID/GID

```
$ go run user_ns.go
```

```
[gophercon]$ whoami  
grafana
```

```
[gophercon]$ groups  
gdm
```


Real Problem



Problem

```
func main() {  
    cmd := makeCmd(exec.Command("/bin/sh "))  
    cmd.SysProcAttr = &syscall.SysProcAttr{  
        Cloneflags: syscall.CLONE_NEWUTS,  
    }  
  
    syscall.Sethostname([]byte("inner-system"))  
    cmd.Run()  
}
```


`/proc/self/exe`





/proc/self/exe

```
$ ls -al /proc/self/exe
```

```
lrwxrwxrwx 1 Dec  7 19:41 /proc/self/exe -> /bin/ls
```



/proc/self/exe

```
$ cat /bin/cat | wc
    154      767   43256
```

```
$ cat /proc/self/exe | wc
    154      767   43256
```



Helper functions

Helper.go

```
func inContainer() bool {  
    i := flag.Bool("inner", false, "child")  
    flag.Parse()  
    return *i  
}
```

```
func main() {  
    if inContainer() {  
        inner()  
    } else {  
        run()  
    }  
}
```

```
func run() error {  
    cmd := makeCmd(exec.Command("/proc/self/exe", "-inner"))  
    cmd.SysProcAttr = &syscall.SysProcAttr{  
        Cloneflags: syscall.CLONE_NEWUTS,  
    }  
    return cmd.Run()  
}  
  
func inner() error {  
    syscall.Sethostname([]byte("inner-system"))  
    return makeCmd(exec.Command("/bin/sh")).Run()  
}
```



UTS namespace


```
$ go build uts.go
```

```
$ sudo ./uts
```

```
[gophercon]$ hostname
```

```
inner-system
```

Picture or it didn't happen



/proc/self/exe

```
19406 pts/1  S    0:00  |  |  |  \_ sudo ./hello
19416 pts/1  Sl   0:00  |  |  |      \_ ./hello
19421 pts/1  Sl   0:00  |  |  |          \_ /proc/self/exe -inner
19426 pts/1  S+   0:00  |  |  |              \_ /bin/sh
```

Reexec

<https://github.com/moby/moby/tree/master/pkg/reexec>

PID namespace



PID namespace

```
func run() error {  
    cmd := makeCmd(exec.Command("/proc/self/exe", "-inner"))  
    cmd.SysProcAttr = &syscall.SysProcAttr{  
        Cloneflags: syscall.CLONE_NEWPID,  
    }  
    return cmd.Run()  
}  
  
func inner() error {  
    fmt.Println("Inner code PID", os.Getpid())  
}
```



PID namespace

```
$ go build main.go
```

```
$ sudo ./main
```

```
Inner code PID 1
```

```
[gophercon]$
```




Problem

[gophercon]\$ ps ax

PID	TTY	STAT	TIME	COMMAND
1	?	Ss	0:03	/sbin/init splash
2	?	S	0:00	[kthreadd]
4	?	S<	0:00	[kworker/0:0H]
6	?	S<	0:00	[mm_percpu_wq]
7	?	S	0:01	[ksoftirqd/0]
8	?	S	1:32	[rcu_sched]
9	?	S	0:00	[rcu_bh]
10	?	S	0:00	[migration/0]

Mnt namespace



Perspective


```
$ sudo ./main
```

Inner code PID 1

```
[gophercon]$ findmnt -o+PROPAGATION
```

TARGET	OPTIONS	
PROPAGATION		
/	rw,noatime,errors=remount-ro,data=ordered	shared
- /dev	rw,nosuid,relatime,size=7622836k,nr_inodes=1905709,mode=755	shared
- /dev/pts	rw,nosuid,noexec,relatime,gid=5,mode=620,ptmxmode=000	shared
- /dev/shm	rw,nosuid,nodev	shared
- /dev/mqueue	rw,relatime	shared
` - /dev/hugepages	rw,relatime,pagesize=2M	shared
- /run	rw,nosuid,noexec,relatime,size=1530088k,mode=755	shared
- /run/lock	rw,nosuid,nodev,noexec,relatime,size=5120k	shared
` - /run/user/1000	rw,nosuid,nodev,relatime,size=1530084k,mode=700,uid=1000,gid=1000	shared
- /sys	rw,nosuid,nodev,noexec,relatime	shared
- /sys/kernel/security	rw,nosuid,nodev,noexec,relatime	shared

mnt problem



Mnt namespace

```
func run() error {  
    cmd := makeCmd(exec.Command("/proc/self/exe", "-inner"))  
    cmd.SysProcAttr = &syscall.SysProcAttr{  
        Cloneflags: syscall.CLONE_NEWNS  
    }  
    return cmd.Run()  
}  
  
func inner() error {  
    return makeCmd(exec.Command("/bin/sh")).Run()  
}
```

```
$ sudo ./main
```

```
Inner code PID 1
```

```
[gophercon]$ findmnt -o+PROPAGATION
```

TARGET	OPTIONS	
PROPAGATION		
/	rw,noatime,errors=remount-ro,data=ordered	private
- /dev	rw,nosuid,relatime,size=7622836k,nr_inodes=1905709,mode=755	private
- /dev/pts	rw,nosuid,noexec,relatime,gid=5,mode=620,ptmxmode=000	private
- /dev/shm	rw,nosuid,nodev	private
- /dev/mqueue	rw,relatime	private
` - /dev/hugepages	rw,relatime,pagesize=2M	private
- /run	rw,nosuid,noexec,relatime,size=1530088k,mode=755	private
- /run/lock	rw,nosuid,nodev,noexec,relatime,size=5120k	private
` - /run/user/1000	rw,nosuid,nodev,relatime,size=1530084k,mode=700,uid=1000,gid=1000	private
- /sys	rw,nosuid,nodev,noexec,relatime	private
- /sys/kernel/security	rw,nosuid,nodev,noexec,relatime	private

mnt problem

Unshare



unshare

allows a process (or thread) to disassociate parts of its execution context that are currently being shared with other processes (or threads). Part of the execution context, such as the mount namespace, is shared implicitly when a new process is created using [fork\(2\)](#) or [vfork\(2\)](#), while other parts, such as virtual memory, may be shared by explicit request when creating a process or thread using [clone\(2\)](#).

The main use of **unshare()** is to allow a process to control its shared execution context without creating a new process.



unshare flags

```
func run() error {  
    cmd := makeCmd(exec.Command("/proc/self/exe", "-inner"))  
    cmd.SysProcAttr = &syscall.SysProcAttr{  
        Cloneflags: syscall.CLONE_NEWNS | syscall.CLONE_NEWPID,  
        Unshareflags: syscall.CLONE_NEWNS,  
    }  
    return cmd.Run()  
}  
  
func inner() error {  
    syscall.Mount("/proc", "/proc", "proc", uintptr(0), "")  
}
```



PID (isolated)

```
$ go build main.go
```

```
$ sudo ./main
```

```
[gophercon]$ ps -aexf
```

PID	TTY	STAT	TIME	COMMAND
1	pts/0	Sl	0:00	/proc/self/exe -inner PS1=[gophercon]\$
6	pts/0	S	0:00	/bin/sh PS1=[meson10]\$
7	pts/0	R+	0:00	_ ps -aexf PS1=[gophercon]\$ PWD=/



Added in go 1.9

<https://github.com/golang/go/issues/19661>

"It turns out that the systemd developers decided to override the kernel's default setting of 'private' to their own default setting of 'shared'. This means that on Linux machines with systemd, the default is shared , while on Linux machines without systemd, the default is private. Essentially, systemd decided to make it so that there is no default that end programs can rely on. All programs must instead mark the root filesystem as private if they want private namespaces, or as shared if they want shared namespaces if they want to work across all Linux distributions. **I'm pretty sure this was done to frustrate as many people as possible.**"



Homework

- IPC
- Net
- CGroup

xps:~\$ whoami

Piyush Verma

Site Reliability Engineering

Trusting Social

Twitter: **meson10**





Thank you.

Credits

- Sagar Rakshe
- Mohan Dutt Parashar
- Talina Shrotriya
- Akshat Goyal